SLR206: Solutions for Quiz 5

1 Uninitialized queues

Recall that if a queue initially stores $\{winner, loser\}$, then two processes can solve consensus by performing a dequeue operation and deciding on your own value if you are the winner, or the value of the other process otherwise.

Suppose that we can only use empty queues. The trick is to use two queues, one for each process. Each process p_i first initializes its queue Q[j], then registers its input in T[i], and then for j=0,1 (in this order) runs the consensus algorithm $Cons_j$ using the initialized queue Q[j] and proposing the value decided in the first consensus to the second one. If for some j=0,1, $T[j]=\bot$ (the input of p_j is not yet registered), p_i simply skips the corresponding consensus (Algorithm 1).

The proof of correcteness is left as an exercise.

Hint: To prove that Algorithm 1 indeed solves consensus, assume that p_i (i = 01,) was the first process to write in T[i]. Show first that if both processes return, then they both go through consensus $Cons_i$ and, thus, they must return the same value (returned by $Cons_i$).

Note that the approach can be used for any set of uninitialized base objects and any algorithm that solves consensus among any number of processes assuming a specific initialization of these base objects.

Algorithm 1 2-process consensus using empty queues

```
1: Shared variables:
 2:
      registers T[0,1] = \{\bot\}
 3:
      queues Q[0,1] = \{\}
 4: propose(v_i) performed by p_i (i = 0, 1):
 5:
      Q[i].eng(winner);
 6:
      Q[i].eng(loser);
 7:
      T[i].write(v_i);
 8:
      v = v_i;
      for j = 0..1 do
 9:
10:
         if T[j] \neq \bot then
           v = Cons_i(v) (using queue Q[j]);
11:
12:
       return(v);
```

2 Lock-free universal construction

If we only lock-freedom is required, processes may avoid sharing their operations with each other. Instead, whenever a process has a new operation op to be performed, it may go to the next (not yet accessed) consensus objects in the series $C[1], C[2], \ldots$, let it be C[k], and propose (op, i) to it by invoking C[k]. propose(op, i). The value (op', j) returned by C[k] is appended to the local list decided. If (op', j) = (op, i), the process returns the response of op computed,

based on *decided*, using the sequential specification of the implemented object. Please check that the algorithm is correct.

3 Consensus numbers of TAS

By contradiciton, suppose that an algorithm A solves binary 3-process consensus (for processes p_0, p_1, p_2) using registers and TAS objects.

Recall that any input configuration C_0 in which some process p proposes 0 and another process q proposes 1 is *bivalent*: p running solo from C_0 must decide 0 and q running solo from C_0 must decide 1.

We show that C_0 must have a critical descendant: a configuration C reachable from C_0 by a finite execution such that:

- C is bivalent:
- for each p_i (i = 0, 1, 2), $C.p_i$ (the configuration obtained after p_i takes one more step of A after C) is monovalent (0-valent or 1-valent).

Indeed, suppose, by contradiction, that C_0 has no critical descendants. Thus, for every bivalent descendant of C_0 (including C_0 itself) has a bivalent descendant.

Now we construct an infinite execution that only goes through bivalent configurations as follows. Let C_1 be the bivalent one-step extension of C_0 (it must exist by our assumption), C_2 - the bivalent one-step extension of C_1 , etc. We denote the resulting infinite execution by E. Recall that no process can decide in a bivalent configuration - otherwise, the agreement property of consensus is violated in some extension of this configuration. Thus, no process can decide in E—a violation of the termination property of consensus.

Thus, A has a critical configuration C. Without loss of generality let $C.p_0$ (the extension of C with one step of p_0) be 0-valent, and $C.p_1$ be 1-valent.

We observe that the steps of p_0 and p_1 enabled in C must be on the same base object X: otherwise they commute, i.e., configurations $C.p_0.p_1$ and $C.p_1.p_0$ are indistinguishable (the process and base-object states are identical), but have opposite valences.

Moreover, as we have shown in the class, X cannot be a register (to see this, consider the cases of read and write operations performed by p_0 and p_1 and show that in each case we can find indistinguishable configurations of opposite valences).

Note that until now we have not used the assumption that A uses only registers and TAS objects. The claims above hold for any wait-free consensus algorithm using base objects.

Thus, X must be a TAS object. But then $C.p_0.p_1$ and $C.p_1.p_0$ only differ in the local states of p_0 and p_1 : only these two processes "know" who won the TAS object and who lost it, and all base objects have identical states in the two configurations.

Thus, p_2 running solo from $C.p_0.p_1$ must decide the same value as it would decide running solo from $C.p_1.p_0$ —a contradiction with the assumption that $C.p_0$ (and, thus, $C.p_0.p_1$) is 0-valent and $C.p_1$ (and, thus, $C.p_1.p_0$) is 1-valent.

Hence, TAS and registers cannot be used to solve consensus among 3 processes, which, combined with the 2-process consensus algorithm using TAS and registers discussed in class, implies that the consensus number of TAS is 2.